

Optimization of a DSL for a Linear Algebra approach to OLAP

[Tech Report](#)

[Presentation 1](#)

[Presentation 2](#)

[Final report](#) 

João Afonso • João Fernandes

2016/2017

13 Feb



◎ Project Introduction

- The challenge
- HPC goal

◎ Work methodology definition

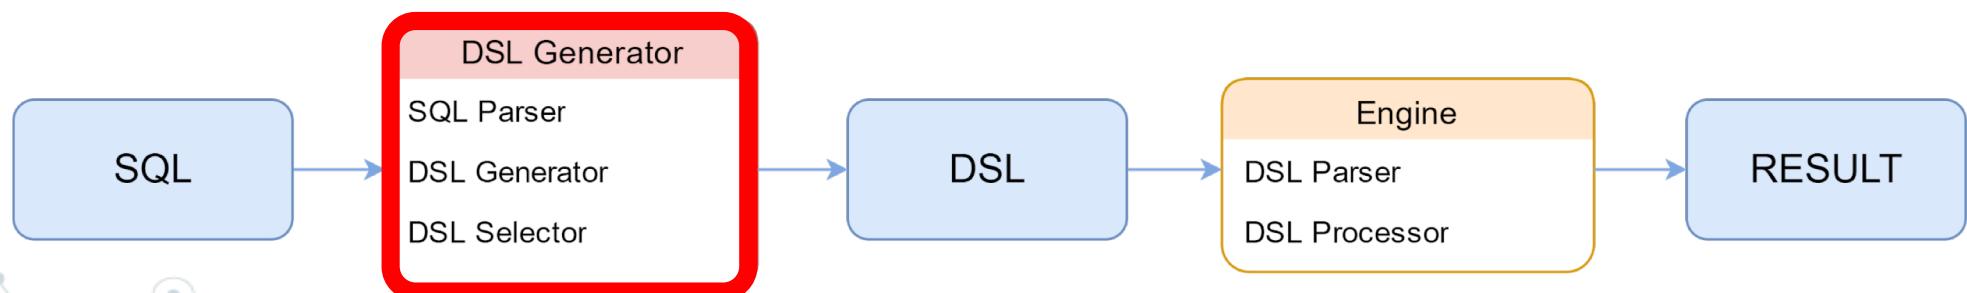
◎ Work for the week

- Read all the given code and documentation



The challenge

- ◎ To automatically translate any SQL query to the already defined DSL
 - For now, limited to SELECT queries



HPC goal

- ◎ Provide the most efficient DSL query for any SQL query and test machine
- ◎ Do it fast!

Work methodology

- ◎ Weekly presentation of the current progress:
 - Every Wednesday at 3pm
 - Professors:
 - Alberto Proença
 - José Nuno Oliveira
 - Collaborators:
 - Filipe Oliveira
 - Rogério Pontes

22 Feb



- Study the given git repositories
- Study the given documentation
- Define a clear path to develop the project
- Determine the necessary resources



Already Implemented

- ◎ A “hard”coded prototype of a simplified TPC-H query
 - Some code can be reused to implement the actual DSL engine;
- ◎ A DSL validator only

Necessary steps

◎ Implement a SQL to DSL converter

- Implement a SQL parser
- Design a data structure that can represent a SQL query
- Load SQL to intermediate data structure
- Design a data structure that can represent the DSL
- Generate all possible conversion paths from SQL to DSL
- Find a way to choose the most efficient one

Necessary resource

- ◎ Decide on a programming language
 - Preferably ANSI C
 - Possibly Java or C++
- ◎ SQL grammar in the chosen programming language

01 Mar

- ◎ Understand the SQL to Diagram conversion process
 - Simple queries validation with professor JNO
- ◎ Definition of the initial steps in SQL processing

Possible mistake?

	c_0	c_1	c_2	c_3	c_4	c_5	c_6	c_7	c_8	c_9
r_0	1	1	1	1	1	1	1	0	0	0
r_1	0	0	0	0	0	0	0	1	1	0
r_2	0	0	0	0	0	0	0	0	0	1

Figure 2: Dense projection matrix produced from the collection of raw data (adapted from TPC-H benchmark lineitem table), from column #9 (return flag column) with 2-way association achieved recurring to GQuarks.

	c_0	c_1	c_2	c_3	c_4	c_5	c_6	c_7	c_8	c_9
r_0	0	0	0	0	0	0	0	0	0	0
r_1	0	0	0	0	0	0	0	0	0	0
r_2	0	0	0	0	0	0	0	0	0	0
r_3	1	1	1	1	1	1	1	0	0	0
r_4	0	0	0	0	0	0	0	1	1	1

Figure 3: Dense projection matrix produced from the collection of raw data (adapted from TPC-H benchmark lineitem table), from column #10 (line status column) with 2-way association achieved recurring to GQuarks.

	c_0	c_1	c_2	c_3	c_4	c_5	c_6	c_7	c_8	c_9
r_0	0	0	0	0	0	0	0	0	0	0
r_1	0	0	0	0	0	0	0	0	0	0
r_2	0	0	0	0	0	0	0	0	0	0
r_3	0	0	0	0	0	0	0	0	0	0
r_4	0	0	0	0	0	0	0	0	0	0
r_5	0	0	0	0	0	0	0	0	0	0
r_6	0	0	0	0	0	0	0	0	0	0
r_7	0	0	0	0	0	0	0	0	0	0
r_8	0	0	0	0	0	0	0	0	0	0
r_9	0	0	0	0	0	0	0	1	1	0
r_{10}	0	0	0	0	0	0	0	0	0	0
r_{11}	0	0	0	0	0	0	0	0	0	0
r_{12}	0	0	0	0	0	0	0	0	0	0
r_{13}	0	0	0	0	0	0	0	0	0	0
r_{14}	0	0	0	0	0	0	0	0	0	1

Figure 4: Projection matrix produced from the Khatri-Rao product between return flag and line status columns from lineitem table, from columns #9 and #10.

Deveria ser:
 $q = m \times o$
 $r = n = p$

Given the matrices A with dimensions $(m \times n)$ and B with dimensions $(o \times p)$, the resulting projection matrix has dimensions $(q \times r)$, being $q = m = o$ and $r = m \times o$. The resulting matrix dimensions can also be expressed as $((m \times o) \times n)$ as shown in figure 4.

SQL to Diagram

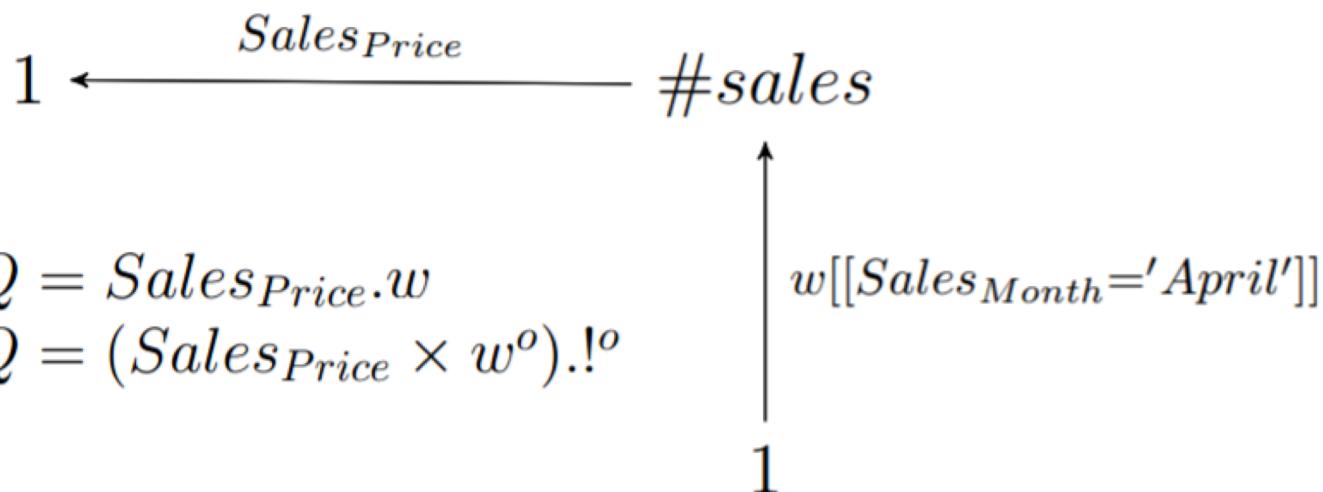
```
SELECT Sales.Price  
FROM Sales
```

1 ← *SalesPrice* #sales

Sales	
Price	Month
10.4	April
11.6	May
10.1	June
15.0	December
16.3	April

SQL to Diagram

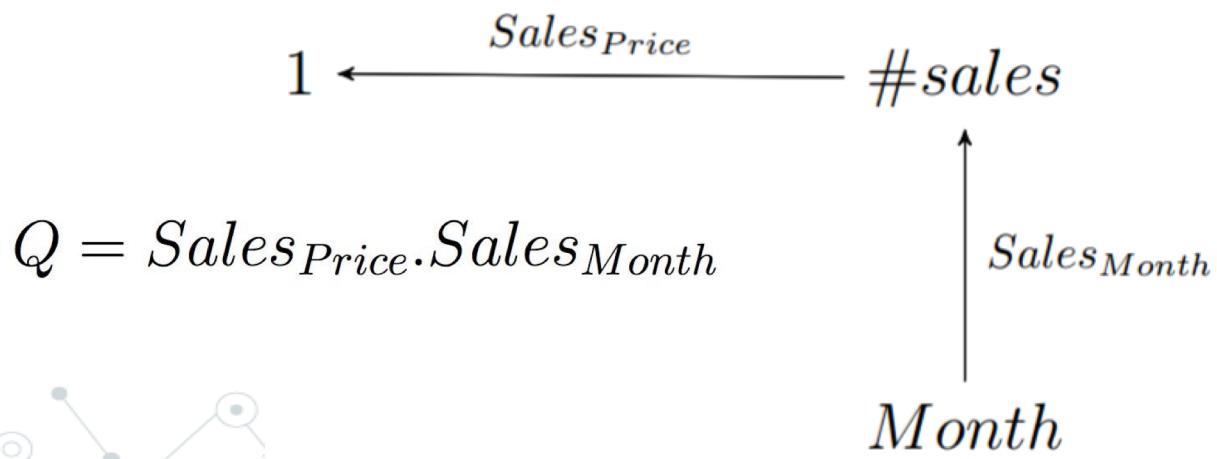
```
SELECT sum(Sales.Price)
FROM Sales
WHERE Sales.Month='April'
```



Sales	
Price	Month
10.4	April
11.6	May
10.1	June
15.0	December
16.3	April

SQL to Diagram

```
SELECT sum(Sales.Price)  
FROM Sales  
GROUP BY Sales.Month
```



Sales	
Price	Month
10.4	April
11.6	May
10.1	June
15.0	December
16.3	April

SQL to Diagram

◎ Aggregation functions

- SUM
- COUNT (Bang Operator)
- AVG (SUM & COUNT)

SQL to LA method



- ◎ Make “bird’s eye” of the query topology
 - Query has to be written in verbose
 - `SELECT A.attr (...)` instead of `SELECT attr (...)`
 - Possibly implemented as a graph (**YES!**)
 - Considering the **WHERE** clauses
- ◎ Distinguish the Measures and Dimensions
 - A Measure has to be aggregated
 - A Dimension has to be a member of **GROUP BY**



SQL to LA method



- ◎ Decide on a output format for the data
 - $A \rightarrow B$ YES during processing
or Maybe $(A \times B) \rightarrow 1$ YES at the end
 - Isomorphic
 - How do we choose between them?
- ◎ Find the most efficient path that *(to study...)*



08 Mar

Professor João Saraiva contribution

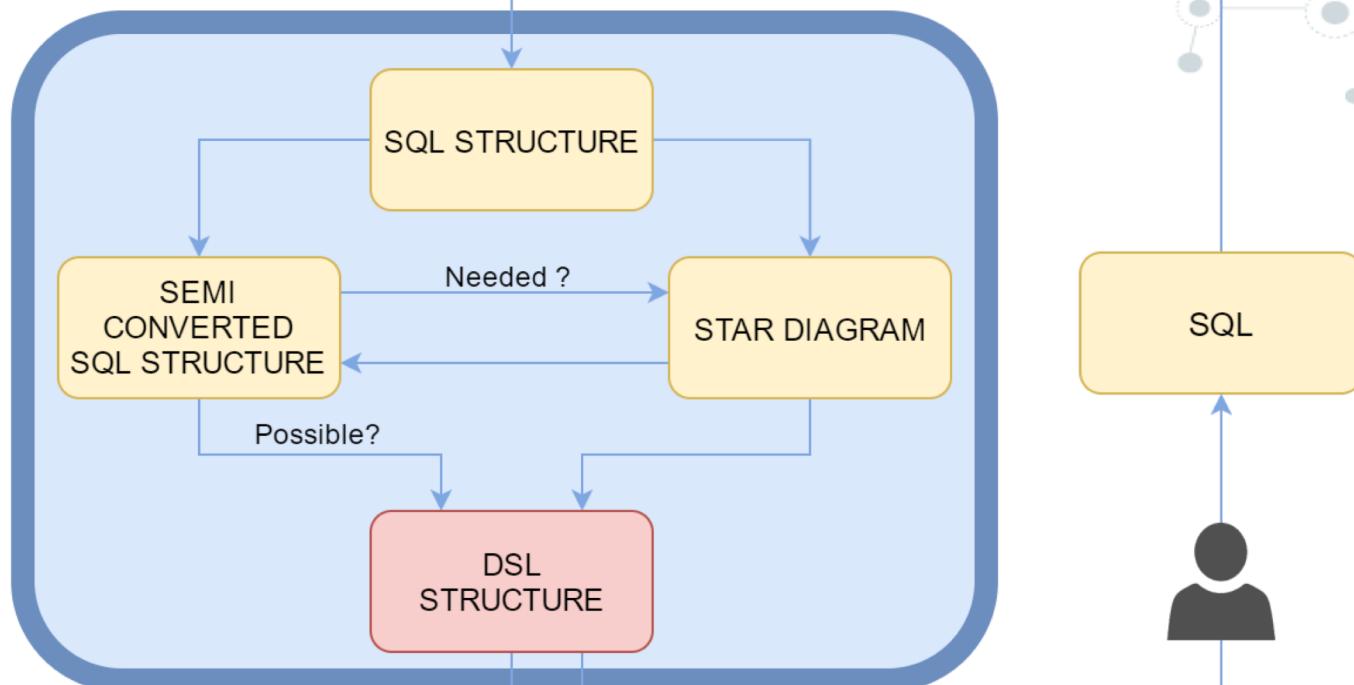
- ◎ Python as prototyping language:
 - C - Slower development, too verbose
 - Java - Too verbose
 - Haskell - Harder to incrementally solve the problem
- ◎ C as implementation language:
 - The disadvantages are irrelevant
 - Higher performance

Python Magic

- ◎ Functional prototype made in 3 days
- ◎ Structure:
 - SQL Parser - Hard coded structure (TODO)
 - In memory SQL to LA - converter.py
 - Engine - engine.py

Current Dilemma

Converter



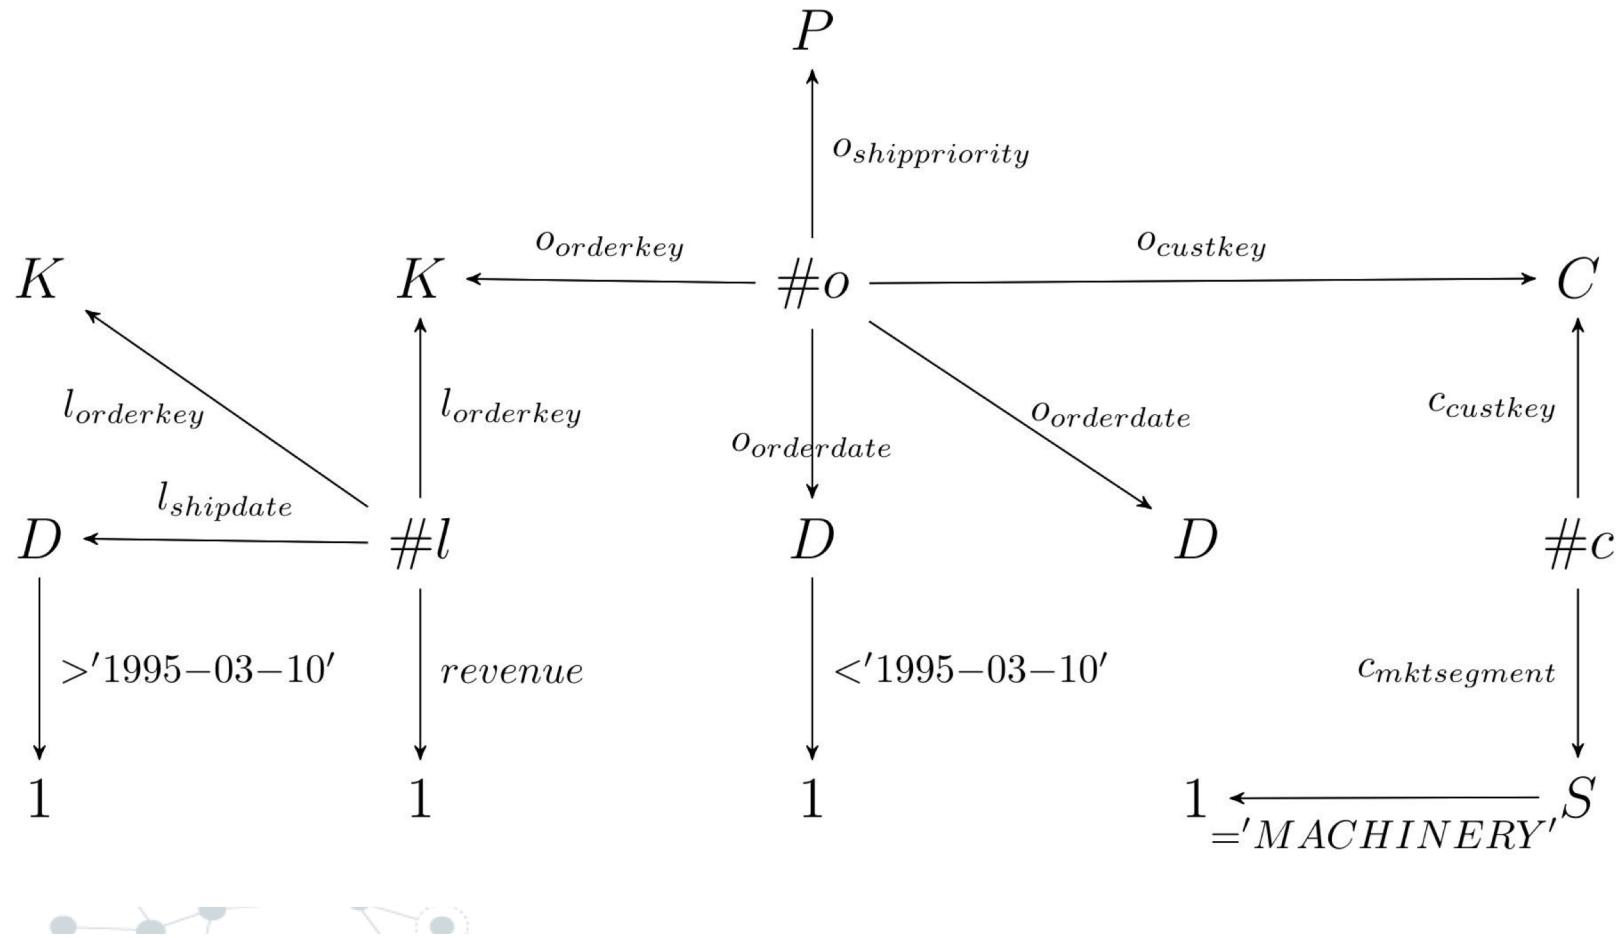
15 Mar



- ◎ Hands on star diagram to LA conversion
 - Can we multiply every matrix with any product (if the types match) without losing data? **NO**
 - If the result matrix dimensions are the same, the output will also be? **NO**



Methodology PDF



Rethink operators

- Conjunctive Hadamard: WHERE AND
- Disjunctive Hadamard: WHERE OR
- Dot product
- Problem JOIN and OR:

```
SELECT A.a, B.b  
FROM A, B  
WHERE A.id = B.b  
OR A.a > 10
```

Rethink operators

- ◎ Khatri-Rao
- ◎ Horizontal Khatri-Rao:

- $1 \rightarrow A \leftarrow \#t \rightarrow A$
 $(M^\circ \cdot N) \nabla N \Leftrightarrow (M^\circ \nabla N^\circ)^\circ$
 - $1 \rightarrow A \leftarrow \#t$
 $(M^\circ \nabla N^\circ)^\circ \Leftrightarrow M \triangleright N$



$$M = \begin{bmatrix} 5 \\ 15 \\ 25 \end{bmatrix} > 10 = \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix}$$

$$Data = [5 \quad 15 \quad 25 \quad 15]$$

$$N = Bitmap(Data) = \begin{bmatrix} 0 & 1 & 2 & 3 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{matrix} 5 \\ 15 \\ 25 \end{matrix}$$

$$M^o = [0 \quad 1 \quad 1]$$

$$M^o.N = \begin{bmatrix} 0 & 1 & 2 & 3 \\ 0 & 1 & 1 & 1 \end{bmatrix}$$

$$(M^o.N) \triangle N = \begin{bmatrix} 0 & 1 & 2 & 3 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{matrix} 5 \\ 15 \\ 25 \end{matrix}$$

$$M = \begin{bmatrix} 5 \\ 15 \\ 25 \end{bmatrix} > 10 = \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix}$$

$$M^o = [0 \ 1 \ 1]$$

$$Data = [5 \ 15 \ 25 \ 15]$$

$$N = Bitmap(Data) = \begin{bmatrix} 0 & 1 & 2 & 3 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{array}{l} 5 \\ 15 \\ 25 \end{array}$$

$$(M^o \triangleright N^o) = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \begin{array}{l} 0 \\ 1 \\ 2 \\ 3 \end{array}$$

$$N^o = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \begin{array}{l} 0 \\ 1 \\ 2 \\ 3 \end{array}$$

$$(M^o \triangleright N^o)^o = \begin{bmatrix} 0 & 1 & 2 & 3 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{array}{l} 5 \\ 15 \\ 25 \end{array}$$

$$Data = [5 \quad 15 \quad 25 \quad 15]$$

$$M = \begin{bmatrix} 5 \\ 15 \\ 25 \end{bmatrix} > 10 = \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix}$$

$$N = Bitmap(Data) = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{matrix} 5 \\ 15 \\ 25 \end{matrix}$$

$$M \triangleright N = \begin{bmatrix} 0 & 1 & 2 & 3 \\ 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{matrix} 5 \\ 15 \\ 25 \end{matrix}$$

+ Performance!

22 Mar

- Dependency Strategy
- Cost Model
- HPC, where is it?
- Filters
- Group By Order
- Starting point in methodology

Dependency Strategy

- ◎ Allow the engine to easily recognize possible operation level parallelism:
 - Eg.: $A \times B \times C \times D$
 - $((A \times B) \times C) \times D$
 - $(A \times B) \times (C \times D)$
- ◎ Should independent instructions be explicit or not:
 - Eg.: PAR { $A \times B$; $C \times D$; }
 - Eg.: $A \times B$; $C \times D$; (only)

Filters

- ◎ What should we do with filters of the same table?
 - Multiple simple filters:
 - **Hadamard**(filter(C1) , filter(C2));
 - One complex filter
 - **filter(C1 and C2);**
 - Could it be optimized with HEP-Frame?
- ◎ The problem here is complex expression solving:
 - (A and B) or C or (D and E)



Cost Model

- ◎ Hadamard
- ◎ Horizontal Khatri-Rao (Vec Mat)
- ◎ Khatri-Rao (Vec Mat)
- ◎ Dot Product
- ◎ Khatri-Rao (Mat Mat)

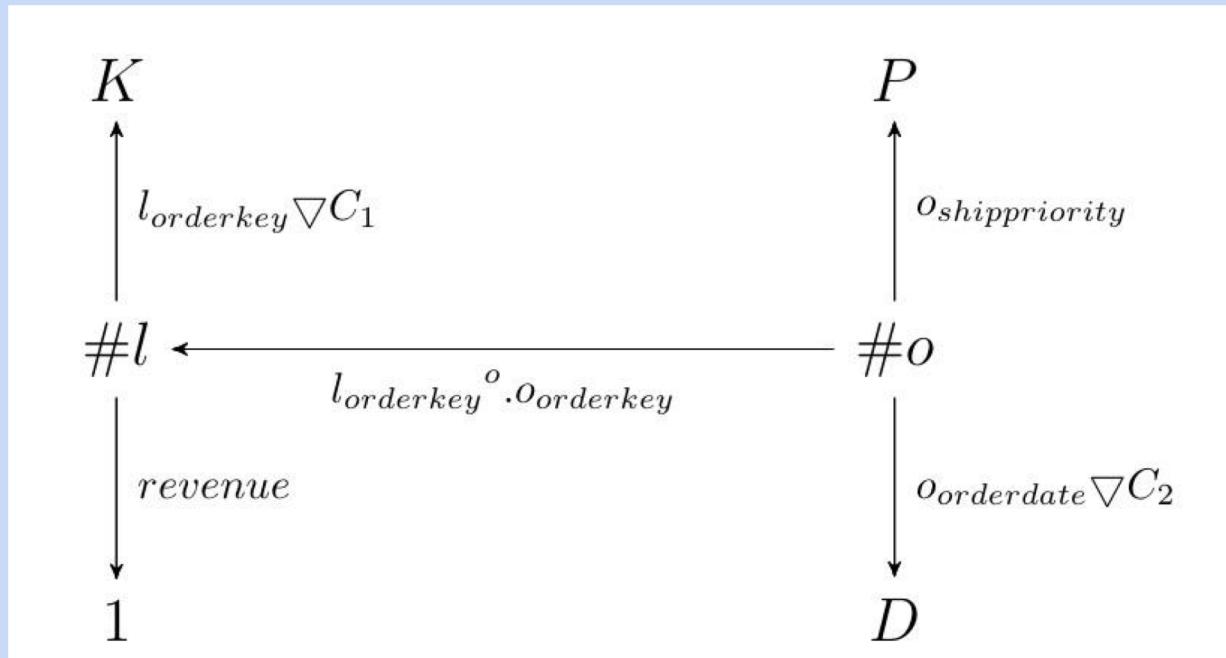


HPC, where is it?

- ◎ Should the end code be ported to C?
 - Currently, it runs instantly
 - It does not benefit from parallelism at all
 - Harder to maintain (in case the project keeps going)
- ◎ The HPC part of the project comes from the output quality

Group By Order

- ◎ TPC-H Query 3:
 - What if the **GROUP BY** order was:
 - $P \times K \times D$
 - Instead of:
 - $P \times D \times K$



Prototype State

1. Solve **WHERE** conditions
 1. **Hadamard** between filters on the same attribute
 2. Apply **Horizontal Khatri-Rao** where the topology allows
 3. **Dot Product** of remaining filters with the corresponding attribute
 1. No more filters
 4. **Khatri-Rao** of conditions with the **GROUP BY** dimensions in the same table
 5. If unique, **Dot Product** of remaining conditions with the table **JOIN**
 6. **Khatri-Rao** of conditions with the first **JOIN** in the table
 1. No more conditions
2. **Khatri-Rao** of same table **GROUP BY** dimensions
3. **Hadamard** of the conditions with the measures of the same table
4. **Dot Product** of the **JOINS** (create both directions)

30 Mar

- ◎ The group cannot be present at the weekly meeting on wednesday 29th



05 Apr

© Presentation



OR (different tables)

Multiple typed diagrams.

Pros:

1. Introduced a new level of abstraction

Cons:

1. Most of the abstracted operations are replicated for each diagram
 1. Major hit in performance

Can we define a methodology to make this the last step (or possibly delete it)?

26 Apr

◎ TODO list:

- Methodology
 - Aggregation functions
 - Subqueries
- Documentation
 - Apply methodology to all TPC-H queries
- Code everything

Methodology (revised)

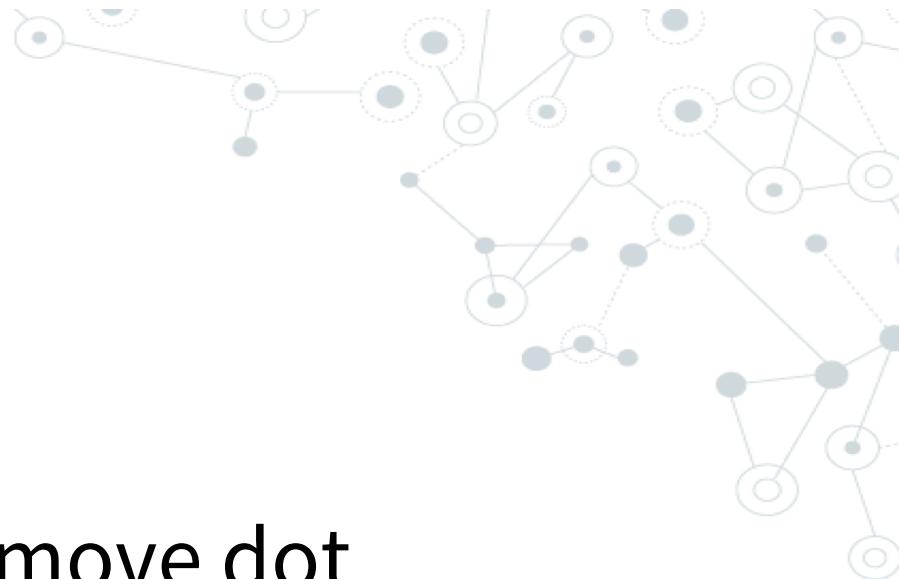
1. Solve in-table operations
 1. Filters
 1. **Hadamard** between filters on the same attribute
 2. Apply **Horizontal Khatri-Rao** where the topology allows
 3. **Dot Product** of remaining filters with the corresponding attribute
 2. Conditions
 1. **Hadamard AND/OR** of all conditions

Methodology (revised)

1. ...
2. Dot Product of the **JOINS** (in the opposite direction to the facts table)
3. Replicate graphs depending on the existence of **ORs**
4. While there is more than one dimension matrix
 1. Merge (Dot Product) **JOINs** with the **GROUP BY** and the conditions
 2. **Khatri-Rao** between all dimensions (**Hadamard** if both are 1)
5. Join the graphs (**Hadamard OR**)
6. Perform the necessary operations depending on the aggregation functions

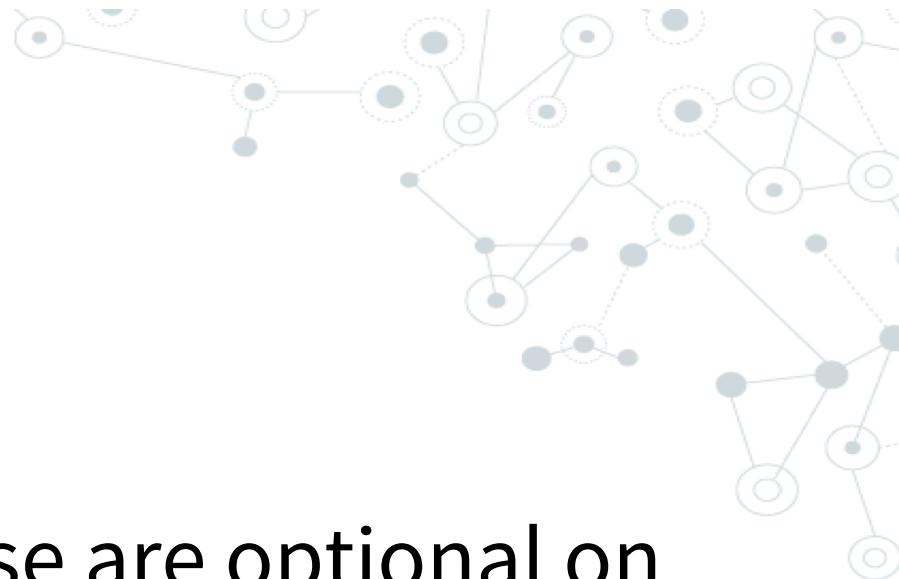
3 May

- ◎ Rethinked operations (remove dot products)
 - JOINS
 - Aggregations



JOINS

- ◎ Dot product and transpose are optional on JOINS:
 - One of the attributes must be the primary key
 - The PK bitmap is a diagonal matrix (identity)
- ◎ $\text{pk} = \text{id}$
- ◎ $\text{id}^o \cdot \text{fk} = \text{fk}$



Agggregation

S

◎ Old version

- Sum: measure . matrix
- Count: bang . matrix
- Avg: Sum/Count

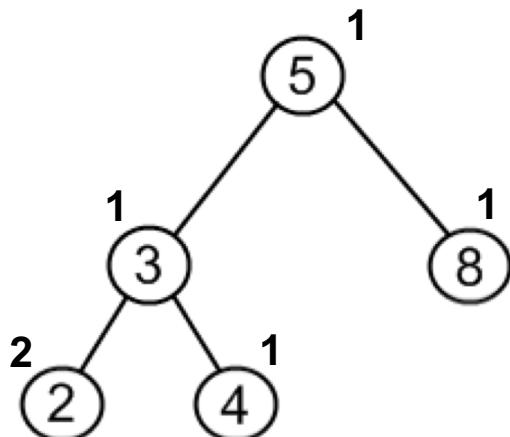
Agggregation

S

◎ New version (eyes on performance)

- Count: Implement the bang.matrix operation based on a search tree
 - CHECK NEXT SLIDE
- Sum: Khatri-Rao then Count
- Average: “Parallel” count and sum

	#0	#1	#2	#3	#4	#5
1: A						
2: B		1		1		
3: C				1		
4: D			1			
5: D						1
6: E						
7: F						
8: G	1					



◎ Query 3

- ~ **3 600 000 000** empty rows
- ~ **0 000 011 000** rows with data

◎ Problem

- `for (i=0; i<n_rows; i++) ...`
- Explained by the other group

Methodology Problem?

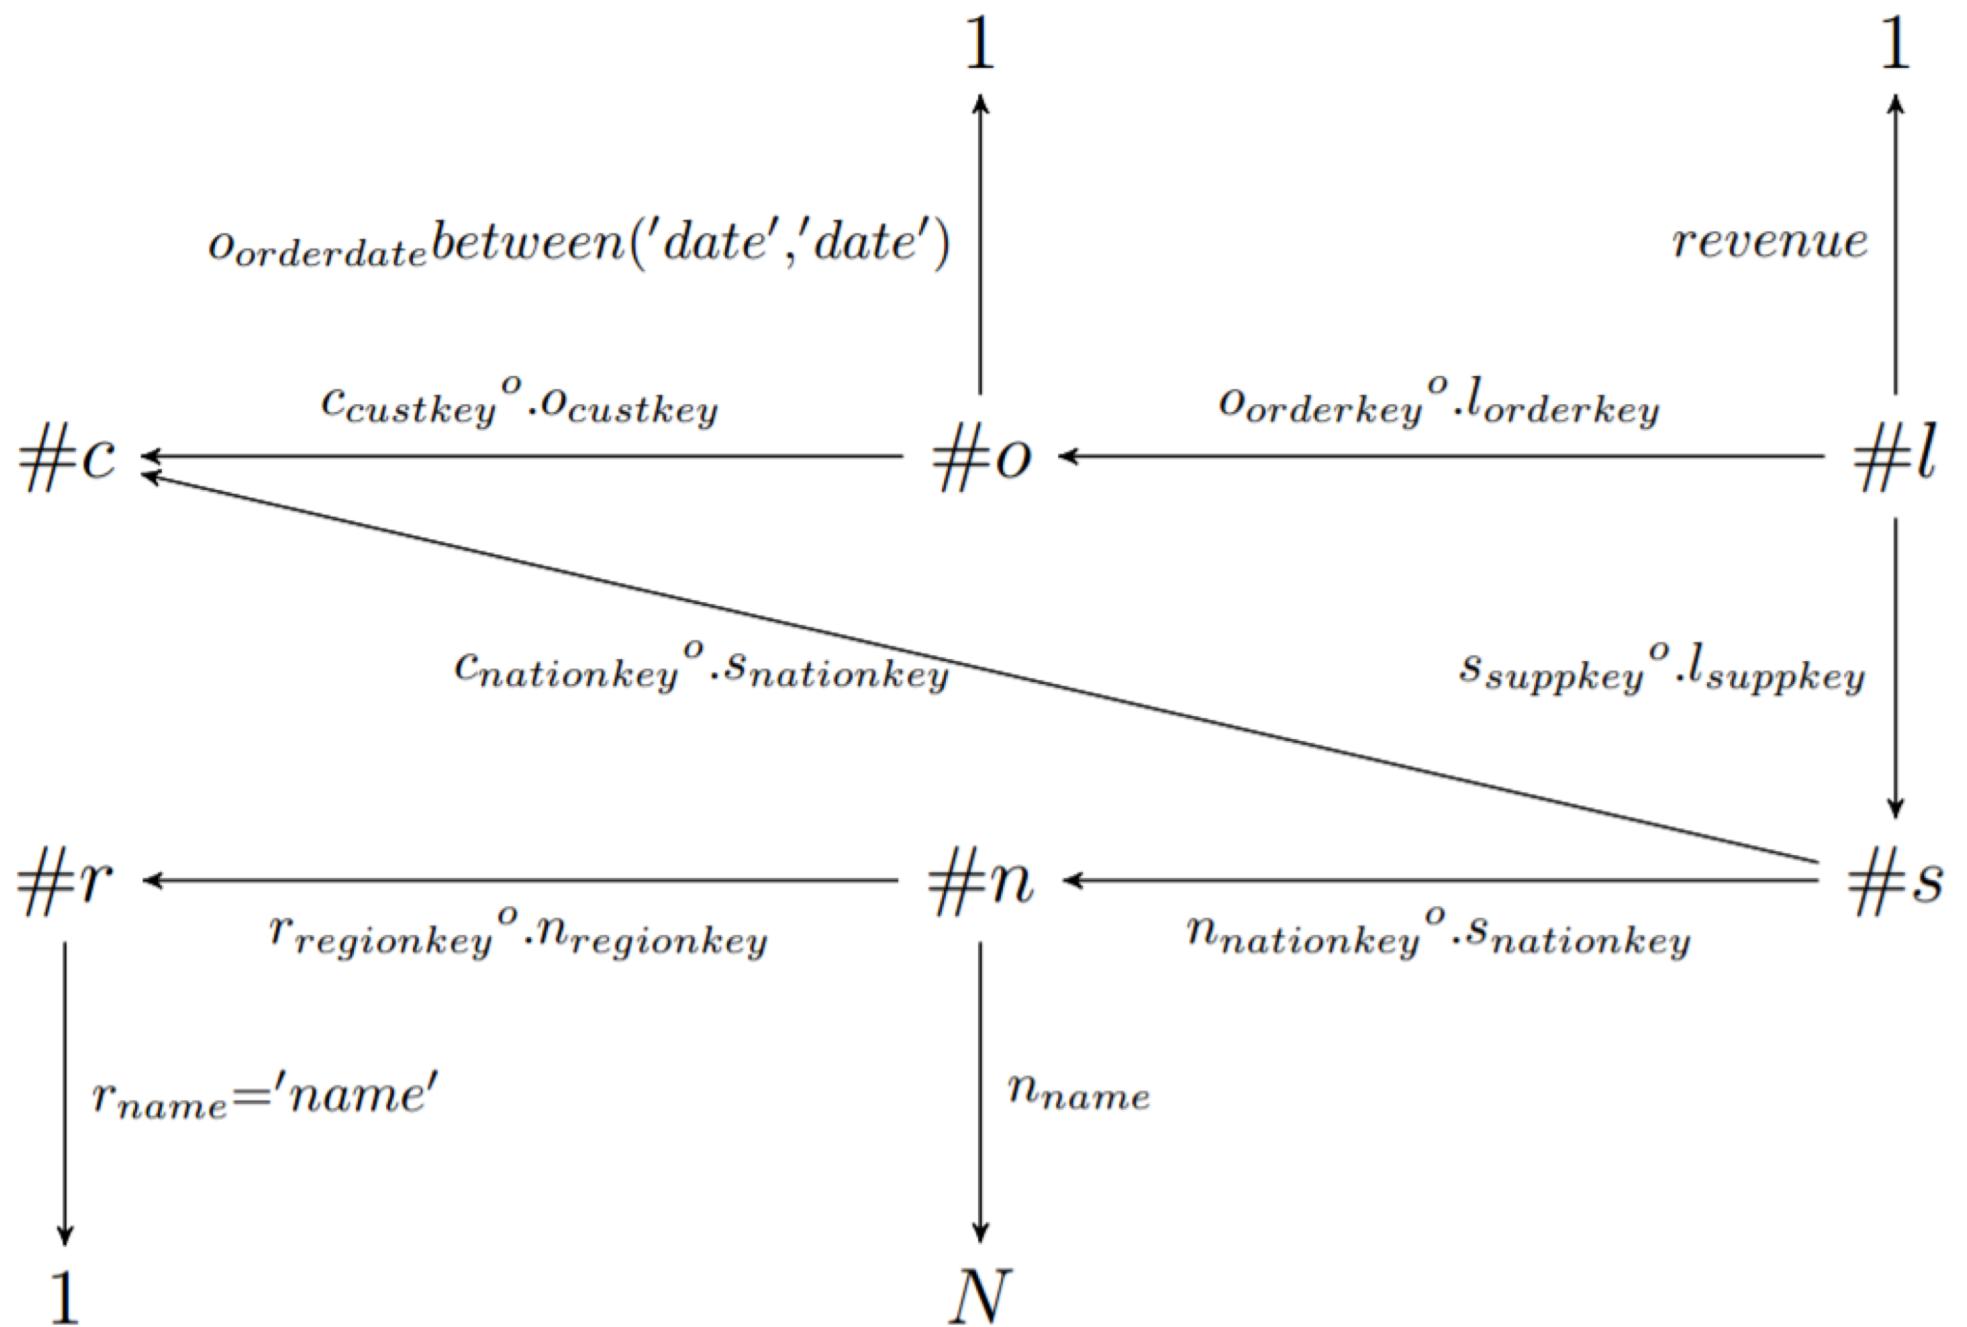


- ◎ TPC-H query 5:
 - What if the graph contains a loop?



Query 5

```
select
    n_name,
    sum(l_extendedprice * (1 - l_discount)) as revenue
from
    customer, orders, lineitem, supplier, nation, region
where
    c_custkey = o_custkey
    and l_orderkey = o_orderkey
    and l_suppkey = s_suppkey
    and c_nationkey = s_nationkey
    and s_nationkey = n_nationkey
    and n_regionkey = r_regionkey
    and r_name = ':1'
    and o_orderdate >= date ':2'
    and o_orderdate < date ':2' + interval '1' year
group by
    n_name
order by
    revenue desc;
```



30 May

- Methodology formalization problems
- Query 17 analysis



Formalization problems

- ◎ Operations in different tables have distinct solutions from the ones in the same table
 - Primary data types are **table_vector** and **table_bitmap**?
- ◎ Some operations use more than one line of SQL code (sometimes separated).
 - Eg. **hadamard** between filters in separated **where** clauses

Query 17 analysis

- ◎ Execute query:

```
select
    sum(l_extendedprice) / 7.0 as avg_yearly
from
    lineitem,
    part
where
    p_partkey = l_partkey
    and p_brand = 'Brand#13'
    and p_container = 'SM JAR'
    and l_quantity < (
        select
            0.2 * avg(l_quantity)
        from
            lineitem
        where
            l_partkey = p_partkey
    );
```

Query 17 analysis

- ◎ Get subquery value:

```
select
    0.2 * avg(l_quantity)
from
    lineitem
where
    l_partkey = p_partkey
```

avg = 5.10...

Query 17 analysis

- ◎ Execute query with replaced value

```
select
    sum(l_extendedprice) / 7.0 as avg_yearly
from
    lineitem,
    part
where
    p_partkey = l_partkey
    and p_brand = 'Brand#13'
    and p_container = 'SM JAR'
    and l_quantity < 5.10...
```

7 June

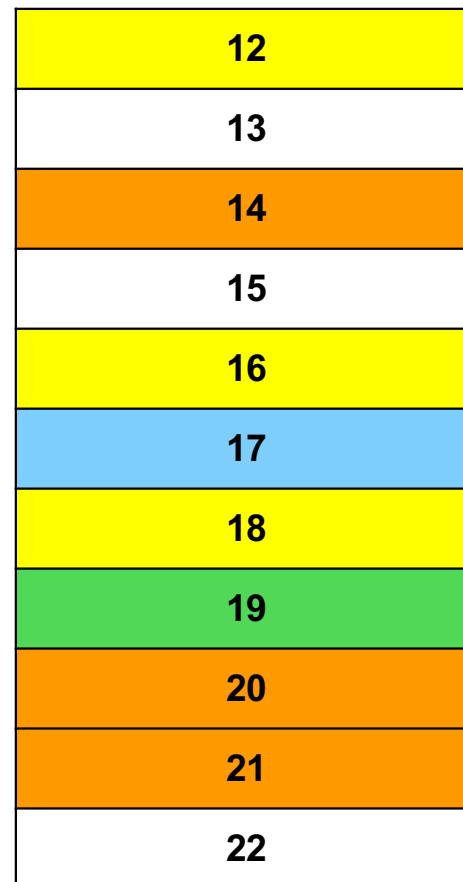
- ◎ Project overview

- Methodology coverage
- Hand solved queries

- ◎ Report structure



Methodology coverage



Methodology problems

1
2
3
4
5
6
7
8
9
10
11

12
13
14
15
16
17
18
19
20
21
22

Not OLAP
Triangular join
Select From Select

Report

◎ Introduction (**1.5 pages**)

- Context
 - OLAP
 - Linear algebra
 - TPC-H
 - Previous work
- Challenges & Goals

Report

◎ Foundations

- Algebraic operations (**2 pages**)
- Data representation (**1 page**)
 - Dense vectors
 - Sparse bitmaps
- Type diagrams (**1 page**)
 - Operations to diagrams
- Simple example (**3 pages**)



Report

◎ Simple example (3 pages)

```
SELECT sum(l_extendedprice)  
FROM lineitem, orders  
WHERE l_orderkey = o_orderkey  
AND o_orderdate > '1996-05-03'  
AND o_totalprice > 100000
```

Report

◎ Design & Implementation

- Performance issues (**1 pages**)
- The approach (**2 pages**)
- Validation (**5 pages: q3**)
- **Coverage**

◎ Conclusion

- Critical analysis (**1 page**)
- Future work (**1 page**)

14 June

- ◎ Solving $\text{Join}^\circ.\text{Join}$
 - For each FK returns all the equal FKS
- ◎ By the id properties:
 - $\text{Join} = p^\circ.f = \text{id}^\circ.f = f$
 - Note: p = primary key; f = foreign key
- ◎ TPC-H query 17 needs this operation:
 - $M.f^\circ.f$
 - Note: the first dot product is an **average** not a **sum**



M. f° . f

- ◎ Using a matrix example we got:
 - $M. f^\circ . f = (((M \nabla f) . !^\circ \nabla f^\circ) . !^\circ)^\circ$
- ◎ It's easier even though it doesn't look like it
- ◎ As explained in slide 46:
 - $(A. !^\circ)^\circ = \text{avl_to_vec}(\text{avl_sum}(A))$
 - Simplified as: $\text{sum}(A)$
 - This operation is identical for: avg, min, max, sum and count.