

Code Parallelizer Goal: to exploit parallel IO and cps processing in both SMP and MPP settings. Rationale: throwing more resources to solve a complex query helps, provided it is easy to determine that parallel processing recovers the administrative overhead Approach: every flow path segment can be handled by an independent process thread. Impact: high

Query Evaluation Maps Goal: to avoid touching any tuple that is not relevant for answering a query. Rationale: the majority of work in solving a query is to disperse tuples of no interest and to find correlated tuples through join conditions. Ideally the database learns these properties over time and re-organizes (or builds a map) to replace disregarding by map lookup. Approach: piggyback selection and joins as database fragmentation instructions Impact: high

MAL Compiler (tactics) Goal: to avoid interpretation of functional expressions Rationale: interpretation of arithmetic expressions with an interpreter is always expensive. Replacing a complex arithmetic expression with a simple dynamically compiled C-functions often pays off. Especially for cached (MAL) queries Approach: Impact: high

Dynamic Query Scheduler (tactical) Goal: to organize the work in a way so as to optimize resource usage Rationale: straight interpretation of a query plan may not lead to the best use of the underlying resources. For example, the content of the runtime cache may provide an opportunity to safe time by accessing a cached source Approach: query scheduling is the last step before a relation algebra interpreter takes over control. The scheduling step involves a re-ordering of the instructions within the boundaries imposed by the flow graph. Impact: medium

Aggregate Groups Goal: to reduce the cost of computing aggregate expressions over time Rationale: many of our applications call for calculation of aggregates over dynamically defined groupings. They call for lengthy scans and it pays to piggyback all aggregate calculates, leaving their result in the cache for later consumption (eg the optimizers) Approach: Impact: High

Demand Driven Interpreter (tactical) Goal: to use the best interpreter and libraries geared at the task at hand Rationale: interpretation of a query plan can be based on different computational models. A demand driven interpretation starts at the intended output and 'walks' backward through the flow graph to collect the pieces, possibly in a pipelined fashion. (Volcano model) Approach: merely calls for a different implementation of the core operators Impact: high

Iterator Strength Reduction Goal: to reduce the cost of iterator execution by moving instructions out of the loop. Rationale: although iteration at the MAL level should be avoided due to the inherent low performance compared to built-in operators, it is not forbidden. In that case we should confine the iterator block to the minimal work needed. Approach: inspect the flowgraph for each iterator and move instructions around. Impact: low

Accumulator Evaluation Goal: to replace operators with cheaper ones. Rationale: based on the actual state of the computation and the richness of the supporting libraries there may exists alternative routes to solve a query. Approach: Operator rewriting depends on properties. No general technique. The first implementation looks at calculator expressions such as they appear frequently in the RAM compiler. Impact: high Pre-requirement: should be called after common term optimization to avoid clashes. Status: Used in the SQL optimizer

Code Inliner Goal: to reduce the calling depth of the interpreter and to obtain a better starting point for code sequencing Rationale: substitution of code blocks (or macro expansion) leads to longer linear code sequences. This provides opportunities for sequencing. Moreover, at runtime building and managing a stackframe is rather expensive. This should be avoided for functions called repeatedly. Impact: medium Status: Used in the SQL optimizer to handle SQL functions.

Code Outliner Goal: to reduce the program size by replacing a group with a single instruction Rationale: inverse macro expansion leads to shorter linear code sequences. This provides opportunities for less interpreter overhead, and to optimize complex, but repetitive instruction sequences with a single hardware call Approach: called explicitly to outline a module (or symbol) Impact: medium

Garbage Collector Goal: to release resources as quickly as possible Rationale: BATs referenced from a MAL program keep resources locked. Approach: In cooperation with a resource scheduler we should identify those that can be released quickly. It requires a forced garbage collection call at the end of the BAT's lifespan. Impact: large Status: Implemented. Algorithm based on end-of-life-span analysis.

Foreign Key replacements Goal: to improve multi-attribute joins over foreign key constraints Rationale: the code produced by the SQL frontend involves foreign key constraints, which provides many opportunities for speedy code using a join index. Impact: large Status: Implemented in the SQL strategic optimizer

Building Blocks

Building blocks mk Tue, 03/30/2010 - 12:07

Optimizer Building Blocks

Some operators are independent of the execution context. In particular, expressions over side-effect free functions with constant parameters could be evaluated before the program block is considered further.

A major task for an optimizer is to select instruction (sequences) which can and should be replaced with cheaper ones. The cost model underlying this decision depends on the processing stage and the overall objective. For example, based on a symbolic analysis their may exist better implementations within the interpreter to perform the job (e.g. hashjoin vs mergejoin). Alternative, expensive intermediates may be cached for later use.

Plan enumeration is often implemented as a Memo structure, which designates alternative sub-plans based on a cost metric. Perhaps we can combine these memo structures into a large table for all possible combinations encountered for a user.

The MAL language does not imply a specific optimizer to be used. Its programs are merely a sequence of specifications, which is interpreted by an engine specific to a given task. Activation of the engine is controlled by a scenario, which currently includes two hooks for optimization; a strategic optimizer and a tactical optimizer. Both engines take a MAL program and produce a (new/modified) MAL program for execution by the lower layers.

MAL programs end-up in the symbol table linked to a user session. An optimizer has the freedom to change the code, provided it is known that the plan derived is invariant to changes in the environment. All others lead to alternative plans, which should be collected as a trail of MAL program blocks. These trails can be inspected for a posteriori analysis, at least in terms of some statistics on the properties of the MAL program structures automatically. Alternatively, the trail may be pruned and reoptimized when appropriate from changes in the environment.

The rule applied for all optimizers is to not return before checking the state of the MAL program, and to assure the dataflow and variable scopes are properly set. It costs some performance, but the difficulties that arise from optimizer interference are very hard to debug. One of the easiest pitfalls is to derive an optimized version of a MAL function while it is already referenced by or when polymorphic typechecking is required afterwards.

Building Your Own Optimizer

Implementation of your own MAL-MAL optimizer can best be started from refinement of one of the examples included in the code base. Beware that only those used in the critical path of SQL execution are thoroughly tested. The others are developed up to the point that the concept and approach can be demonstrated.

The general structure of most optimizers is to actively copy a MAL block into a new program structure. At each step we determine the action taken, e.g. replace the instruction or inject instructions to achieve the desired goal.

A tally on major events should be retained, because it gives valuable insight in the effectiveness of your optimizer. The effects of all optimizers is collected in a system catalog.

Each optimizer ends with a strong define line, `define new(s)`. It performs a complete type and data flow analysis before returning. Moreover, if you are in debug mode, it will keep a copy of the plan produced for inspection. Studying the differences between optimizer steps provide valuable information to improve your code.

The functionality of the optimizer should be clearly delineated. The guiding policy is that it is always safe to not apply an optimizer step. This helps to keep the optimizers as independent as possible.

It really helps if you start with a few tiny examples to test your optimizer. They should be added to the Tests directory and administered in TestSuite.

Breaking up into different components and grouping them together in arbitrary sequences calls for careful programming.

One of the major hurdles is to test interference of the optimizer. The test set is a good starting point, but does not guarantee that all cases have been covered.

In principle, any subset of optimizers should work flawlessly. With a few tens of optimizers this amounts to potential millions of runs. Adherence to a partial order reduces the problem, but still is likely to be too resource consumptive to test continuously.

Lifespans

Lifespans mk Tue, 03/30/2010 - 12:09

Optimizers may be interested in the characteristic of the barrier blocks for making a decision. The variables have a lifespan in the code blocks, denoted by properties `beginLifespan`, `endLifespan`. The `beginLifespan` denotes the instruction where it receives its first value, the `endLifespan` the last instruction in which it was used as operand or target.

If, however, the last use lies within a BARRIER block, we can not be sure about its end of life status, because a block redo may implicitly revive it. For these situations we associate the `endLifespan` with the block exit.

In many cases, we have to determine if the lifespan interferes with a optimization decision being prepared. The lifespan is calculated once at the beginning of the optimizer sequence. It should either be maintained to reflect the most accurate situation while optimizing the code base. In particular, it means that any move/remove/addition of a MAL instruction calls for either a recalculation or further propagation. Similar what will be the best strategy for the time being we just recode.

See if all arguments mentioned in the instruction at point pc are still visible at instruction pc and have not been updated in the mean time. Take into account that variables may be declared inside a block. This can be calculated using the BARRIER/CATCH and EXIT pairs.

The safety property should be relatively easy to determine for each MAL function. This calls for accessing the function MAL block and to inspect the arguments of the signature.

Any instruction may block identification of a common subexpression. It suffices to stumble upon an unsafe function whose parameter lists has a non-empty intersection with the targeted instruction. To illustrate, consider the sequence

```
13 -> f(A,B,C);
14 -> g(A,B,F);
15 -> g(A,B,C);
16 -> h()
```

The instruction G1—g(A,B,C) is blocking if G1 is an alias for (A,B,C). Alternatively, function g() may be unsafe and (D,E,F) has a non-empty intersection with (A,B,C). An alias can only be used later on for readability (and not be used for a function with side effects).

Alias removal

Alias removal mk Tue, 03/30/2010 - 16:01

The routine `optimizer.aliasRemoval()` walks through the program looking for simple assignment statements, e.g. `V=W`. It replaces all subsequent occurrences of V by W provided V is assigned a value once and W does not change in the remainder of the code. Special care should be taken for iterator blocks as illustrated in the case below:

```
1--0;
2--"name";
3--"name";
4--"name";
5--"name";
6--"name";
7--"name";
8--"name";
9--"name";
10--"name";
11--"name";
12--"name";
13--"name";
14--"name";
15--"name";
16--"name";
17--"name";
18--"name";
19--"name";
20--"name";
21--"name";
22--"name";
23--"name";
24--"name";
25--"name";
26--"name";
27--"name";
28--"name";
29--"name";
30--"name";
31--"name";
32--"name";
33--"name";
34--"name";
35--"name";
36--"name";
37--"name";
38--"name";
39--"name";
40--"name";
41--"name";
42--"name";
43--"name";
44--"name";
45--"name";
46--"name";
47--"name";
48--"name";
49--"name";
50--"name";
51--"name";
52--"name";
53--"name";
54--"name";
55--"name";
56--"name";
57--"name";
58--"name";
59--"name";
60--"name";
61--"name";
62--"name";
63--"name";
64--"name";
65--"name";
66--"name";
67--"name";
68--"name";
69--"name";
70--"name";
71--"name";
72--"name";
73--"name";
74--"name";
75--"name";
76--"name";
77--"name";
78--"name";
79--"name";
80--"name";
81--"name";
82--"name";
83--"name";
84--"name";
85--"name";
86--"name";
87--"name";
88--"name";
89--"name";
90--"name";
91--"name";
92--"name";
93--"name";
94--"name";
95--"name";
96--"name";
97--"name";
98--"name";
99--"name";
100--"name";
101--"name";
102--"name";
103--"name";
104--"name";
105--"name";
106--"name";
107--"name";
108--"name";
109--"name";
110--"name";
111--"name";
112--"name";
113--"name";
114--"name";
115--"name";
116--"name";
117--"name";
118--"name";
119--"name";
120--"name";
121--"name";
122--"name";
123--"name";
124--"name";
125--"name";
126--"name";
127--"name";
128--"name";
129--"name";
130--"name";
131--"name";
132--"name";
133--"name";
134--"name";
135--"name";
136--"name";
137--"name";
138--"name";
139--"name";
140--"name";
141--"name";
142--"name";
143--"name";
144--"name";
145--"name";
146--"name";
147--"name";
148--"name";
149--"name";
150--"name";
151--"name";
152--"name";
153--"name";
154--"name";
155--"name";
156--"name";
157--"name";
158--"name";
159--"name";
160--"name";
161--"name";
162--"name";
163--"name";
164--"name";
165--"name";
166--"name";
167--"name";
168--"name";
169--"name";
170--"name";
171--"name";
172--"name";
173--"name";
174--"name";
175--"name";
176--"name";
177--"name";
178--"name";
179--"name";
180--"name";
181--"name";
182--"name";
183--"name";
184--"name";
185--"name";
186--"name";
187--"name";
188--"name";
189--"name";
190--"name";
191--"name";
192--"name";
193--"name";
194--"name";
195--"name";
196--"name";
197--"name";
198--"name";
199--"name";
200--"name";
201--"name";
202--"name";
203--"name";
204--"name";
205--"name";
206--"name";
207--"name";
208--"name";
209--"name";
210--"name";
211--"name";
212--"name";
213--"name";
214--"name";
215--"name";
216--"name";
217--"name";
218--"name";
219--"name";
220--"name";
221--"name";
222--"name";
223--"name";
224--"name";
225--"name";
226--"name";
227--"name";
228--"name";
229--"name";
230--"name";
231--"name";
232--"name";
233--"name";
234--"name";
235--"name";
236--"name";
237--"name";
238--"name";
239--"name";
240--"name";
241--"name";
242--"name";
243--"name";
244--"name";
245--"name";
246--"name";
247--"name";
248--"name";
249--"name";
250--"name";
251--"name";
252--"name";
253--"name";
254--"name";
255--"name";
256--"name";
257--"name";
258--"name";
259--"name";
260--"name";
261--"name";
262--"name";
263--"name";
264--"name";
265--"name";
266--"name";
267--"name";
268--"name";
269--"name";
270--"name";
271--"name";
272--"name";
273--"name";
274--"name";
275--"name";
276--"name";
277--"name";
278--"name";
279--"name";
280--"name";
281--"name";
282--"name";
283--"name";
284--"name";
285--"name";
286--"name";
287--"name";
288--"name";
289--"name";
290--"name";
291--"name";
292--"name";
293--"name";
294--"name";
295--"name";
296--"name";
297--"name";
298--"name";
299--"name";
300--"name";
301--"name";
302--"name";
303--"name";
304--"name";
305--"name";
306--"name";
307--"name";
308--"name";
309--"name";
310--"name";
311--"name";
312--"name";
313--"name";
314--"name";
315--"name";
316--"name";
317--"name";
318--"name";
319--"name";
320--"name";
321--"name";
322--"name";
323--"name";
324--"name";
325--"name";
326--"name";
327--"name";
328--"name";
329--"name";
330--"name";
331--"name";
332--"name";
333--"name";
334--"name";
335--"name";
336--"name";
337--"name";
338--"name";
339--"name";
340--"name";
341--"name";
342--"name";
343--"name";
344--"name";
345--"name";
346--"name";
347--"name";
348--"name";
349--"name";
350--"name";
351--"name";
352--"name";
353--"name";
354--"name";
355--"name";
356--"name";
357--"name";
358--"name";
359--"name";
360--"name";
361--"name";
362--"name";
363--"name";
364--"name";
365--"name";
366--"name";
367--"name";
368--"name";
369--"name";
370--"name";
371--"name";
372--"name";
373--"name";
374--"name";
375--"name";
376--"name";
377--"name";
378--"name";
379--"name";
380--"name";
381--"name";
382--"name";
383--"name";
384--"name";
385--"name";
386--"name";
387--"name";
388--"name";
389--"name";
390--"name";
391--"name";
392--"name";
393--"name";
394--"name";
395--"name";
396--"name";
397--"name";
398--"name";
399--"name";
400--"name";
401--"name";
402--"name";
403--"name";
404--"name";
405--"name";
406--"name";
407--"name";
408--"name";
409--"name";
410--"name";
411--"name";
412--"name";
413--"name";
414--"name";
415--"name";
416--"name";
417--"name";
418--"name";
419--"name";
420--"name";
421--"name";
422--"name";
423--"name";
424--"name";
425--"name";
426--"name";
427--"name";
428--"name";
429--"name";
430--"name";
431--"name";
432--"name";
433--"name";
434--"name";
435--"name";
436--"name";
437--"name";
438--"name";
439--"name";
440--"name";
441--"name";
442--"name";
443--"name";
444--"name";
445--"name";
446--"name";
447--"name";
448--"name";
449--"name";
450--"name";
451--"name";
452--"name";
453--"name";
454--"name";
455--"name";
456--"name";
457--"name";
458--"name";
459--"name";
460--"name";
461--"name";
462--"name";
463--"name";
464--"name";
465--"name";
466--"name";
467--"name";
468--"name";
469--"name";
470--"name";
471--"name";
472--"name";
473--"name";
474--"name";
475--"name";
476--"name";
477--"name";
478--"name";
479--"name";
480--"name";
481--"name";
482--"name";
483--"name";
484--"name";
485--"name";
486--"name";
487--"name";
488--"name";
489--"name";
490--"name";
491--"name";
492--"name";
493--"name";
494--"name";
495--"name";
496--"name";
497--"name";
498--"name";
499--"name";
500--"name";
501--"name";
502--"name";
503--"name";
504--"name";
505--"name";
506--"name";
507--"name";
508--"name";
509--"name";
510--"name";
511--"name";
512--"name";
513--"name";
514--"name";
515--"name";
516--"name";
517--"name";
518--"name";
519--"name";
520--"name";
521--"name";
522--"name";
523--"name";
524--"name";
525--"name";
526--"name";
527--"name";
528--"name";
529--"name";
530--"name";
531--"name";
532--"name";
533--"name";
534--"name";
535--"name";
536--"name";
537--"name";
538--"name";
539--"name";
540--"name";
541--"name";
542--"name";
543--"name";
544--"name";
545--"name";
546--"name";
547--"name";
548--"name";
549--"name";
550--"name";
551--"name";
552--"name";
553--"name";
554--"name";
555--"name";
556--"name";
557--"name";
558--"name";
559--"name";
560--"name";
561--"name";
562--"name";
563--"name";
564--"name";
565--"name";
566--"name";
567--"name";
568--"name";
569--"name";
570--"name";
571--"name";
572--"name";
573--"name";
574--"name";
575--"name";
576--"name";
577--"name";
578--"name";
579--"name";
580--"name";
581--"name";
582--"name";
583--"name";
584--"name";
585--"name";
586--"name";
587--"name";
588--"name";
589--"name";
590--"name";
591--"name";
592--"name";
593--"name";
594--"name";
595--"name";
596--"name";
597--"name";
598--"name";
599--"name";
600--"name";
601--"name";
602--"name";
603--"name";
604--"name";
605--"name";
606--"name";
607--"name";
608--"name";
609--"name";
610--"name";
611--"name";
612--"name";
613--"name";
614--"name";
615--"name";
616--"name";
617--"name";
618--"name";
619--"name";
620--"name";
621--"name";
622--"name";
623--"name";
624--"name";
625--"name";
626--"name";
627--"name";
628--"name";
629--"name";
630--"name";
631--"name";
632--"name";
633--"name";
634--"name";
635--"name";
636--"name";
637--"name";
638--"name";
639--"name";
640--"name";
641--"name";
642--"name";
643--"name";
644--"name";
645--"name";
646--"name";
647--"name";
648--"name";
649--"name";
650--"name";
651--"name";
652--"name";
653--"name";
654--"name";
655--"name";
656--"name";
657--"name";
658--"name";
659--"name";
660--"name";
661--"name";
662--"name";
663--"name";
664--"name";
665--"name";
666--"name";
667--"name";
668--"name";
669--"name";
670--"name";
671--"name";
672--"name";
673--"name";
674--"name";
675--"name";
676--"name";
677--"name";
678--"name";
679--"name";
680--"name";
681--"name";
682--"name";
683--"name";
684--"name";
685--"name";
686--"name";
687--"name";
688--"name";
689--"name";
690--"name";
691--"name";
692--"name";
693--"name";
694--"name";
695--"name";
696--"name";
697--"name";
698--"name";
699--"name";
700--"name";
701--"name";
702--"name";
703--"name";
704--"name";
705--"name";
706--"name";
707--"name";
708--"name";
709--"name";
710--"name";
711--"name";
712--"name";
713--"name";
714--"name";
715--"name";
716--"name";
717--"name";
718--"name";
719--"name";
720--"name";
721--"name";
722--"name";
723--"name";
724--"name";
725--"name";
726--"name";
727--"name";
728--"name";
729--"name";
730--"name";
731--"name";
732--"name";
733--"name";
734--"name";
735--"name";
736--"name";
737--"name";
738--"name";
739--"name";
740--"name";
741--"name";
742--"name";
743--"name";
744--"name";
745--"name";
746--"name";
747--"name";
748--"name";
749--"name";
750--"name";
751--"name";
752--"name";
753--"name";
754--"name";
755--"name";
756--"name";
757--"name";
758--"name";
759--"name";
760--"name";
761--"name";
762--"name";
763--"name";
764--"name";
765--"name";
766--"name";
767--"name";
768--"name";
769--"name";
770--"name";
771--"name";
772--"name";
773--"name";
774--"name";
775--"name";
776--"name";
777--"name";
778--"name";
779--"name";
780--"name";
781--"name";
782--"name";
783--"name";
784--"name";
785--"name";
786--"name";
787--"name";
788--"name";
789--"name";
790--"name";
791--"name";
792--"name";
793--"name";
794--"name";
795--"name";
796--"name";
797--"name";
798--"name";
799--"name";
800--"name";
801--"name";
802--"name";
803--"name";
804--"name";
805--"name";
806--"name";
807--"name";
808--"name";
809--"name";
810--"name";
811--"name";
812--"name";
813--"name";
814--"name";
815--"name";
816--"name";
817--"name";
818--"name";
819--"name";
820--"name";
821--"name";
822--"name";
823--"name";
824--"name";
825--"name";
826--"name";
827--"name";
828--"name";
829--"name";
830--"name";
831--"name";
832--"name";
833--"name";
834--"name";
835--"name";
836--"name";
837--"name";
838--"name";
839--"name";
840--"name";
841--"name";
842--"name";
843--"name";
844--"name";
845--"name";
846--"name";
847--"name";
848--"name";
849--"name";
850--"name";
851--"name";
852--"name";
853--"name";
854--"name";
855--"name";
856--"name";
857--"name";
858--"name";
859--"name";
860--"name";
861--"name";
862--"name";
863--"name";
864--"name";
865--"name";
866--"name";
867--"name";
868--"name";
869--"name";
870--"name";
871--"name";
872--"name";
873--"name";
874--"name";
875--"name";
876--"name";
877--"name";
878--"name";
879--"name";
880--"name";
881--"name";
882--"name";
883--"name";
884--"name";
885--"name";
886--"name";
887--"name";
888--"name";
889--"name";
890--"name";
891--"name";
892--"name";
893--"name";
894--"name";
895--"name";
896--"name";
897--"name";
898--"name";
899--"name";
900--"name";
901--"name";
902--"name";
903--"name";
904--"name";
905--"name";
906--"name";
907--"name";
908--"name";
909--"name";
910--"name";
911--"name";
912--"name";
913--"name";
914--"name";
915--"name";
916--"name";
917--"name";
918--"name";
919--"name";
920--"name";
921--"name";
922--"name";
923--"name";
924--"name";
925--"name";
926--"name";
927--"name";
928--"name";
929--"name";
930--"name";
931--"name";
932--"name";
933--"name";
934--"name";
935--"name";
936--"name";
937--"name";
938--"name";
939--"name";
940--"name";
941--"name";
942--"name";
943--"name";
944--"name";
945--"name";
946--"name";
947--"name";
948--"name";
949--"name";
950--"name";
951--"name";
952--"name";
953--"name";
954--"name";
955--"name";
956--"name";
957--"name";
958--"name";
959--"name";
960--"name";
961--"name";
962--"name";
963--"name";
964--"name";
965--"name";
966--"name";
967--"name";
968--"name";
969--"name";
970--"name";
971--"name";
972--"name";
973--"name";
974--"name";
975--"name";
976--"name";
977--"name";
978--"name";
979--"name";
980--"name";
981--"name";
982--"name";
983--"name";
984--"name";
985--"name";
986--"name";
987--"name";
988--"name";
989--"name";
990--"name";
991--"name";
992--"name";
993--"name";
994--"name";
995--"name";
996--"name";
997--"name";
998--"name";
999--"name";
1000--"name";
```

The constant strings are propagated to the

print()

routine, while the initial assignment

is=0

should be retained. The code block becomes:

```
1--0;
2--"name";
3--"name";
4--"name";
5--"name";
6--"name";
7--"name";
8--"name";
9--"name";
10--"name";
11--"name";
12--"name";
13--"name";
14--"name";
15--"name";
16--"name";
17--"name";
18--"name";
19--"name";
20--"name";
21--"name";
22--"name";
23--"name";
24--"name";
25--"name";
26--"name";
27--"name";
28--"name";
29--"name";
30--"name";
31--"name";
32--"name";
33--"name";
34--"name";
35--"name";
36--"name";
37--"name";
38--"name";
39--"name";
40--"name";
41--"name";
42--"name";
43--"name";
44--"name";
45--"name";
46--"name";
47--"name";
48--"name";
49--"name";
50--"name";
51--"name";
52--"name";
53--"name";
54--"name";
55--"name";
56--"name";
57--"name";
58--"name";
59--"name";
60--"name";
61--"name";
62--"name";
63--"name";
64--"name";
65--"name";
66--"name";
67--"name";
68--"name";
69--"name";
70--"name";
71--"name";
72--"name";
73--"name";
74--"name";
75--"name";
76--"name";
77--"name";
78--"name";
79--"name";
80--"name";
81--"name";
82--"name";
83--"name";
84--"name";
85--"name";
86--"name";
87--"name";
88--"name";
89--"name";
90--"name";
91--"name";
92--"name";
93--"name";
94--"name";
```


8.18 Performance profiler

A key issue in developing fast programs using the Monet database back-end requires a keen eye on where performance is lost. Although performance tracking and measurements are highly application dependent, a simple to use tool makes life a lot easier (See [dataflow.org](#) and [imageperf.it](#)).
Activation of the performance monitor has a global effect, i.e. all concurrent actions on the kernel are traced, but the events are only sent to the client who initiated the profiler thread.
The profiler event can be handled in several ways. The default strategy is to ship the event record immediately over a stream to a performance monitor. An alternative strategy is preparation of off-line performance analysis, where event information is retained in tables.
To reduce the interference of performance measurement with the experiments, the user can use an event cache, which is emptied explicitly upon need.

```
module profiler;
pattern activateStream(str,...):void
pattern deactivateStream(str,...):void
pattern queryStream():void
  address CMEQueryProfilerStream comment "Send the events to output stream";
pattern queryStreamDone(str):void
  comment (closeStream):void
  address CMEQueryProfilerStream comment "Stop sending the event records";
pattern deactivate():void
  pattern deactivateQuery():void
  address CMEQueryProfilerStream comment "Generate an event record for every instruction where it is used."
  pattern deactivateRead(str,...):void
  pattern deactivateWrite(str,...):void
  pattern deactivateTime(str,...):void
  pattern deactivateSystemTime(str,...):void
  pattern deactivate():void
  address CMEStartProfiler comment "Start performance tracing";
comment stop():void
  pattern stop():void
  address CMEStopProfiler comment "Stop performance tracing";
comment reset():void
  address CMEClearTrace comment "Clear the profiler traces";
comment start():void
  comment getTrace(str):void
  address CMEGetTrace comment "Get the trace details of a specific event";
pattern getTrace(i):void
  comment (close):void
  address CMECleanup comment "Remove the temporary tables for profiling";
comment getReads():void
  address CMEGetReads comment "Obtain the number of physical reads";
comment getWrites():void
  address CMEGetWrites comment "Obtain the number of physical reads";
comment getTime():void
  address CMEGetTime comment "Obtain the user timing information.";
comment getSystemTime():void
  address CMEGetSystemTime comment "Obtain the user timing information.";
pattern getTrace():void
```

Remote

Remote mk Tue, 03/30/2010 - 16:34

8.20 Remote querying functionality

Communication with other servers at the MAM level is a delicate task. However, it is indispensable for any distributed functionality. This module provides an abstract way to store and retrieve objects on a remote site. Additionally, functions on a remote site can be executed using objects available in the remote session context. This yields in four primitive functions that form the basis for distribution methods: get, put, register and exec.
The get method simply retrieves a copy of a remote object. Objects can be simple values, strings or BATs. The same holds for the put method, but the other way around. A local object can be stored on a remote site. Upon a successful store, the put method returns the remote identifier for the stored object. With this identifier the object can be addressed, e.g. using the get method to retrieve the object that was stored using put.
The get and put methods are asymmetric. Performing a get on an identifier that was returned by put, results in an object with the same value and type as the one that was put. The result of such an operation is equivalent to making an (expensive) copy of the original object.
The register function takes a local MAM function and makes it known at a remote site. It ensures that it does not overload an already known operation remotely, which could create a semantic conflict. Deregister a function is forbidden, because it would allow for taking over the remote site completely. C-implemented functions, such as io:print() cannot be remotely stored. It would require even more complicated byte[] code shipping and remote compilation to make it work. Currently, the remote procedure may only return a single value.
The choice to let exec only execute functions was made to avoid problems to decide what should be returned to the caller. With a function it is clear and simple to return that what the function signature prescribes. Any side effect (e.g. io:print calls) may cause havoc in the system, but are currently ignored.
This leads to the final contract of this module. The methods should be used correctly, by obeying their contract. Failing to do so will result in errors and possibly undefined behaviour.
The resolve() function can be used to query Merovingian. It returns one or more databases discovered in its vicinity matching the given pattern.

module remote;

module loading and unloading funcs

```
comment activate():void
  address RMTInit comment "Initialise the remote module.";
comment deactivate():void
  address RMTDeinit comment "Release the resources held by the remote module.";
# global connection resolve function
comment resolve(pattern:str):bat():void
  address RMTResolve comment "resolve a pattern against Merovingian and return the URIs";
# session local connection instantiation functions
comment connect(url:str, user:str, passwd:str):str
  address RMTConnect comment "return a newly created connection for url, using user name and password";
comment connect(url:str, user:str, passwd:str, scenario:str):str
  address RMTConnectScen comment "return a newly created connection for url, using user name, password and scenario";
comment disconnect(url:str):void
  address RMTDisconnect comment "disconnects the connection pointed to by handle (received from a call to connect)";
# core transfer functions
pattern get(str:str, ident:str):any
  address RMTGet comment "retrieve a copy of remote object ident";
pattern put(str:str, object:any):str
  address RMTPut comment "copies object to the remote site and returns its identifier";
pattern register(str:str, module:str, fun:str):void
  address RMTRegister comment "register <mod>-<fun> at the remote site";
pattern exec(str:str, func:str):str
  address RMTExec comment "remotely execute <mod>-<fun> and returns the handle to its result";
pattern exec(str:str, module:str, func:str):str
  address RMTExec module comment "remotely execute <mod>-<fun> and returns the handle to its result";
pattern exec(str:str, module:str, fun:str, args:...):str
  address RMTExec module comment "remotely execute <mod>-<fun> using the argument list of remote objects and returns the handle to its result";
pattern exec(str:str, module:str, fun:str, args:...):str
  address RMTExec module comment "remotely execute <mod>-<fun> using the argument list of remote objects and returns the handle to its result";
```

Transactions

Transactions mk Tue, 03/30/2010 - 16:36

8.23 Transaction management

In the philosophy of Monet, transaction management overhead should only be paid when necessary. Transaction management is for this purpose implemented as a module. This code base is largely absolute and should be re-considered when serious OLTP is being supported. Note, however, the SQL front-end obeys transaction semantics.

```
module transaction;
comment open():bat
  address TMOpenSys comment "Save all persistent BATs";
comment commit():bat
  address TMCCommit comment "Global commit on all BATs";
comment abort():bat
  address TMOpenSys comment "Global abort on all BATs";
comment commit(bat:bat,seq,...):bat
  address TMCCommit comment "Commit only a set of BATs, passed in the tail (to which you must have exclusive access)";
pattern commit(seq,...):void
  address TMCCommit comment "Commit changes in certain BATs";
pattern abort(seq,...):void
  address TMAbort comment "Abort changes in certain BATs";
pattern flush(seq,...):void
  address TMAbort comment "Flush a BAT clean without flushing to disk";
comment exec(bat:bat,seq,...):bat
  address TMEExec comment "The previous state of this BAT";
comment exec(bat:bat,seq,...):bat
  address TMEExec comment "Last insertions since last commit";
comment exec(bat:bat,seq,...):bat
  address TMEExec comment "Last deletions since last commit";
```